

Efficient algorithms and data structures

Gregory Kucherov G.Kucherov@skoltech.ru

TA: Stanislav Protasov
Programming exercises:
Ilya Vorobyov I.Vorobyov@skoltech.ru

Course

- ▶ **What the course is:**
 - ▶ a selection of topics on the design and analysis of algorithms
 - ▶ with emphasis on rigorous analysis (Ph.Flajolet: "mathematically oriented engineering")
 - ▶ dealing with basic data structures (graphs, strings, trees, tables, ...)
 - ▶ including programming assignments and in-class projects
- ▶ **What the course is not:**
 - ▶ a programming course
 - ▶ a course oriented to a specific programming language (an imperative programming language is assumed, one of Python, C, C++, Java)
 - ▶ a course oriented to a specific application area
 - ▶ a math course


Course

- ▶ **Varying level of difficulty**
- ▶ **Prerequisites:**
 - ▶ imperative programming (C, C++, Java, ...)
 - ▶ Basic data structures: lists, arrays, stacks, queues
 - ▶ Recursion, Big-Oh notation
 - ▶ Sorting, ...
- ▶ "Free-style" pseudo-code
- ▶ Having a laptop assumed

Grading

- ▶ participation in class 10%
 - ▶ full attendance is expected
 - ▶ in-class projects
- ▶ programming exercises 40%
 - ▶ one every ~2 weeks
 - ▶ plagiarism is not tolerated
- ▶ exam 50%

Useful books



CLRS = Cormen & Leiserson & Rivest & Stein

Some other good algorithm textbooks:

- Steven Skiena, The Algorithm Design Manual, 2nd Edition, Springer, 2008 [a bit advanced?]
- Jon Kleinberg and Éva Tardos, Algorithm Design, MIT Press 2005
- Robert Sedgewick and Kevin Wayne, Algorithms, Addison-Wesley, 4th Edition, 2011 [for beginners, Java-oriented]

How to measure the efficiency of algorithms?


- ▶ **Efficiency (in this course) = TIME and SPACE**
 - ▶ other possible measure of efficiency: accuracy
- ▶ **Classical model: RAM model of computation**
 - ▶ all memory accesses have equal cost
 - ▶ no parallel execution
 - ▶ unit cost ($O(1)$) of basic operations (unless we want to explicitly count individual bits operations)
 - ▶ space = # of computer words (unless bit complexity is considered); each computer word contains $\Theta(\log n)$ bits
 - ▶ other possible measures can be considered: disk accesses, cache misses, probe model, query complexity ...

How to measure the efficiency of algorithms?

- ▶ Algorithms solve mass problems
 - ▶ n : input size (in computer words or bits)
 - ▶ time/space as a function of n
- ▶ Different complexity analyses:
 - ▶ worst-case complexity
 - ▶ average-case complexity
 - ▶ smoothed analysis
 - ▶ query (probe) complexity
 - ▶ ...

Graphs

Directed graph $G = (V, E)$
 V finite set of nodes (vertices)
 $E \subseteq V \times V$ set of edges (arcs),
 i.e., a relation on V




$V = \{1, 2, 3, 4, 5\}$
 $E = \{(1, 2), (1, 3), (2, 3), (3, 2), (4, 4), (4, 5)\}$

Undirected graph $G = (V, E)$
 E set of edges (arcs),
 symmetric relation




$E = \{1, 2, 3, 4\}$
 $V = \{\{1, 2\}, \{1, 3\}, \{1, 4\}, \{2, 4\}\}$


Graphs are everywhere




Acyclic graph of an expression (DAG)




Gene regulation network in biology



Protein-protein interaction network (in yeast)



Social networks



Graph representations

$$G = (V, E) \quad V = \{1, 2, \dots, n\}$$

Adjacency list

reduces the size if $|E| \ll (|V|)^2$
reading time : $O(|V| + |E|)$

Adjacency matrix

using matrix operations
reading time $O(|V|)^2$



Other representations possible

Adjacency lists




$$V = \{1, 2, 3\}$$


$$E = \{(1,1), (1,2), (1,3), (2,1), (3,2)\}$$



weight: $w : A \rightarrow X$



Adjacency matrix




$M[i, j] = 1$ iff j is adjacent to i

$$V = \{1, 2, 3\}$$

$$E = \{(1,1), (1,2), (1,3), (2,1), (3,2)\}$$

$$M = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{pmatrix}$$



weight: $w : A \rightarrow X$

$$W = \begin{pmatrix} 2 & 8 & 5 \\ 3 & 0 & 0 \\ 0 & 2 & 0 \end{pmatrix}$$

Graph algorithms

Exploration

- Depth-first or breadth-first search
- Topological sorting
- Strongly connected components

Path computation

- Shortest path
- Transitive closure
- Eulerian and Hamiltonian paths

Minimum spanning trees

- Kruskal's and Prim's algorithms

Networks

- Maximum flow

Others


- Graph coloring
- Planarity testing
- ...

Shortest paths in graphs

Single-source shortest path: unweighted case

- ▶ Path length = number of edges
- ▶ Distance between two nodes = length of the shortest path
- ▶ *Problem:* given a (directed or undirected) graph $G = (V, E)$ and a source node $s \in V$, compute the distance from s to each reachable node

Single-source shortest path: unweighted case




a subgraph (29,160 nodes) of the graph of Rubik's mini cube (2x2x2) configurations (3,674,160 nodes)

<https://miscellaneouscoder.wordpress.com/2014/07/28/working-with-rubiks-group-cycle-graphs/>


Breadth-first search (BFS)

- Given a source node s ,
- ▶ "Discovers" all nodes reachable from s




Breadth-first search (BFS)

- Given a source node s ,
- ▶ "Discovers" all nodes reachable from s
 - ▶ Proceeds by "concentric circles"
 - ▶ Discovers all nodes at distance d from s before discovering any nodes at distance $d+1$




Breadth-first search (BFS)

- Given a source node s ,
- ▶ "Discovers" all nodes reachable from s
 - ▶ Proceeds by "concentric circles"
 - ▶ Discovers all nodes at distance d from s before discovering any nodes at distance $d+1$
 - ▶ Computes the distances from s
 - ▶ Computes a *breadth-first tree* encoding one shortest path for each node



Breadth-first search (BFS)

- Given a source node s ,
- ▶ "Discovers" all nodes reachable from s
 - ▶ Proceeds by "concentric circles"
 - ▶ Discovers all nodes at distance d from s before discovering any nodes at distance $d+1$
 - ▶ Computes the distances from s
 - ▶ Computes a *breadth-first tree* encoding one shortest path for each node




How it works?

- ▶ colors every node *white* (not yet discovered), *yellow* (discovered but may have white adjacent nodes), or *red* (discovered and all adjacent nodes discovered)
- ▶ yellow nodes = "active frontier" (nodes under processing)
- ▶ when processing a (yellow) node, determine all white neighbors, set their distance to be larger by 1, color them yellow. After that, color the node red.

Breadth-first search (BFS)


```

procedure BFT ( $s$  node of  $V$ ) ;
begin
  for each node  $v$  of  $V$  do {
     $\text{visited}[v] = \text{false}$  ; //  $s$  is white
     $d[v] = \infty$  ;  $\pi(v) = \text{nil}$ 
  }
   $\text{visited}[s] = \text{true}$  ; //  $s$  becomes yellow
   $d[s] = 0$  ;
   $\text{Queue} = \text{enqueue}(\text{empty-queue}, s)$  ;
  while not empty ( $\text{Queue}$ ) do {
     $u = \text{dequeue}(\text{Queue})$  ;
    for  $t = \text{first to last successor of } u$  do
      if not  $\text{visited}[t]$  then
         $\text{visited}[t] = \text{true}$  ; //  $t$  becomes yellow
         $d[t] = d[u] + 1$  ;  $\pi(t) = u$ 
         $\text{Queue} = \text{enqueue}(\text{Queue}, t)$  ;
    //  $s'$  becomes red
  }
end
  
```



wikipedia.org


BFS: example



Queue : a

Order of traversal:


BFS: example



Queue : a b c d e

Order of traversal: a


BFS: example



Queue : a b c d e f g

Order of traversal: a b


BFS: example



Queue : a b c d e f g

Order of traversal: a b c d e

BFS: example



Queue : a b c d e f g

Order of traversal: a b c d e f g

Questions

- ▶ Show that BFS runs in time $O(n+m)$ (assuming the graph is represented by adjacency lists), $n=|V|$, $m=|E|$
- ▶ Show that if (v_1, v_2, \dots, v_i) is the state of the Queue, then $d[v_i] \leq d[v_{i-1}] + 1$ and $d[v_i] \leq d[v_{i+1}]$ for all i
- ▶ Show that upon termination $d[v] = \delta(s, v)$, where $\delta(s, v)$ is the length of the shortest path from s to v

$d[v] = \delta(s, v)$: sketch of the proof

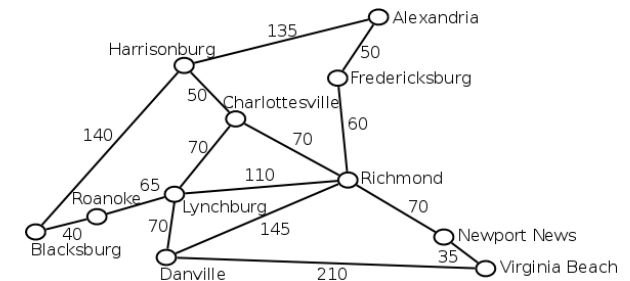
- ▶ by contradiction, let v be the closest to s node with $d[v] > \delta(s, v)$
- ▶ consider a shortest path from s to v , and let u be the node preceding v in this path
- ▶ $\delta(s, v) = \delta(s, u) + 1$ (by properties of shortest paths)
- ▶ consider the moment when u was dequeued ($d[u] = \delta(s, u)$)
- ▶ if v was white then, we have $d[v] = \delta(s, v) \Rightarrow \text{contradiction}$
- ▶ if v was yellow then, it was visited earlier by exploring the successors of some w with $d[w] \leq d[u]$. Then $d[v] = d[w] + 1 \leq d[u] + 1 \Rightarrow \text{contradiction}$
- ▶ if v was red, then $d[v] \leq d[u] \Rightarrow \text{contradiction}$

Space efficient BFS

- ▶ BFS stores the queue which (in the worst case) can contain $O(n)$ nodes, i.e. $O(n \log n)$ bits
- ▶ Can we implement BFS with $o(n \log n)$ bits?
- ▶ **Example of a result:** There exists an algorithm that outputs vertices in the BFS order in time $O(n+m)$ and uses $2n + o(n)$ bits
[N. Banerjee, S. Chakraborty, V. Raman, and S. R. Satti. Space efficient linear time algorithms for BFS, DFS and applications. Theory of Computing Systems, Jan 2018]

Single-source shortest path: weighted case

Single-source shortest path: weighted case



Shortest path problem


Weighted (directed or undirected) graph: $G = (V, E, w)$ where $w : E \rightarrow \mathbf{R}$ (weight/cost)

Source : $s \in V$

Problem: for all $t \in V$, compute

$$\delta(s, t) = \min \{ \{ w(c) ; c \text{ path from } s \text{ to } t \} \cup \{+\infty\} \}$$


Example:



Properties of the shortest paths

Proposition 1 (existence):

shortest paths are well-defined (i.e. for all $t \in V$, $\delta(s, t) > -\infty$) iff the graph does not have a cycle of cost < 0 reachable from s



Proposition 2: if there exists a shortest path from s to t , then there exists one without a cycle

Proposition 3: if there exists a shortest path from s to t , then there exists one with no more than $|V|-1$ edges

Main properties


Property 1: $G = (V, E, w)$

let c be a **shortest** path from p to r and q be the node preceding r in c . Then $\delta(p, r) = \delta(p, q) + w(q, r)$.



Property 2: A subpath of a shortest path is a shortest path

Shortest path tree: tree rooted at s representing shortest paths



Main properties (cont)

Property 3: $G = (V, E, w)$ let c be a path from p to r and q be the node preceding r in c . Then $\delta(p, r) \leq \delta(p, q) + w(q, r)$.



Relaxation

Compute $\delta(s, t)$ by successive approximations

$t \in V$ $d[t]$ = estimate (from above) of $\delta(s, t)$


$\pi[t]$ = predecessor of t on

a path from s to t of cost $d[t]$

Initialization of d and π

INIT

for all $t \in V$ do
 { $d[t] = \infty$; $\pi[t] = \text{nil}$ }
 $d[s] = 0$;



Relaxation of the edge (q, r)

RELAX(q, r)

if $d[q] + w(q, r) < d[r]$
 then { $d[r] = d[q] + w(q, r)$; $\pi[r] = q$ }

Relaxation (cont)

Proposition :

the following property is an invariant of **relax**: for all $t \in V$, $d(t) \geq \delta(s, t)$

Proof: by induction on the number of executions of **relax**

Dijkstra's algorithm




Assumption: $w(p, q) \geq 0$ for all edges (p, q)


```
begin
  INIT;
  S = ∅ ; Q = V ;
  while Q ≠ ∅ do {
    q = MIN_d(Q) ; Q = Q \ {q} ; S = S ∪ {q} ;
    for all r successor of q do
      RELAX(q, r) ;
  }
end
```

- At each iteration, the algorithm extracts a node from Q that is never returned to Q
- RELAX**(q, r) may change $d[r]$

Example




$S = \{s\}$
 $Q = \{a, b, c, d\}$
 $\pi[s] = \text{nil}$
 $\pi[a] = \text{nil}$
 $\pi[b] = \text{nil}$
 $\pi[c] = \text{nil}$
 $\pi[d] = \text{nil}$




$S = \{s, a\}$
 $Q = \{b, c, d\}$
 $\pi[s] = \text{nil}$
 $\pi[a] = s$
 $\pi[b] = \text{nil}$
 $\pi[c] = s$
 $\pi[d] = \text{nil}$

Example (cont)




$S = \{s, a, b\}$
 $Q = \{c, d\}$
 $\pi[s] = \text{nil}$
 $\pi[a] = s$
 $\pi[b] = \text{nil}$
 $\pi[c] = s$
 $\pi[d] = \text{nil}$




$S = \{s, a, b, c\}$
 $Q = \{d\}$
 $\pi[s] = \text{nil}$
 $\pi[a] = s$
 $\pi[b] = a$
 $\pi[c] = a$
 $\pi[d] = \text{nil}$

Example (cont)




$S = \{s, a\}$
 $Q = \{b, c, d\}$
 $\pi[s] = \text{nil}$
 $\pi[a] = s$
 $\pi[b] = a$
 $\pi[c] = a$
 $\pi[d] = \text{nil}$




$S = \{s, a, c\}$
 $Q = \{b, d\}$
 $\pi[s] = \text{nil}$
 $\pi[a] = s$
 $\pi[b] = a$
 $\pi[c] = a$
 $\pi[d] = c$

Example (cont)




$S = \{s, a, c\}$
 $Q = \{b, d\}$, $Q = \{d\}$ then $Q = \emptyset$
 $\pi[s] = \text{nil}$
 $\pi[a] = s$
 $\pi[b] = a$
 $\pi[c] = a$
 $\pi[d] = c$



Correctness of Dijkstra's algorithm

Proposition : After the execution of Dijkstra's algorithm on a graph $G = (V, E, w)$, $d[t] = \delta(s, t)$ for all $t \in V$.


Proof by contradiction: let $d[t] \neq \delta(s, t)$



Properties of Dijkstra's algorithm

Algorithm maintains three sets:

- ▶ S : finished nodes, for which $d[t] = \delta(s, t)$ (red)
- ▶ S' : nodes of Q with $d[t] < \infty$ (yellow)
- ▶ nodes of Q with $d[t] = \infty$ (white)



- ▶ Algorithm can be seen as expanding a ball centered at s following a greedy strategy

Implementation

With adjacency matrix

time $O(n^2)$ (where $n=|V|$)

With adjacency lists

depends on the data structure for Q

we need to support operations:

- insert an element to Q
- extract an element with minimum d value
- modify (decrease) the d value of an element (when relaxing)

\Rightarrow (min-)priority queue

Priority Queues

(max-)Priority Queue is a data structure that supports operations

- ▶ INSERT(S, x)
- ▶ MAX(S)
- ▶ EXTRACT-MAX(S)
- ▶ INCREASE-KEY(S, x, k): increase the key of x to k

Priority Queues are used in

- ▶ Dijkstra's algorithm for shortest paths
- ▶ Prim's algorithm for minimum spanning tree
- ▶ other greedy algorithms


Implemented using heaps

Binary Heaps


Binary heap:

- ▶ a **binary tree** that is
- ▶ **complete**: every level except possibly the bottom one is completely filled and the leaves in the bottom level are as far left as possible
- ▶ satisfies the **(max-)heap property**: the key stored in every node is greater than or equal to the keys stored in its children
 - ▶ If the key at each node is smaller than or equal to the keys of its children, then we have a **min-heap**

Binary (max-)heap: example



Binary heaps stored in arrays



Due to their regular structure, binary heaps are easily stored in arrays

Given index i of a node,

- the index of its parent is $\lfloor i / 2 \rfloor$
- the indices of its children are $2i$ and $2i+1$

18	12	9	8	10	7	5	1	4	3	
1	2	3	4	5	6	7	8	9	10	

element key $A[i]$

index i

Binary heaps: some properties


- ▶ The height of a heap is $\lfloor \log(n) \rfloor$
- ▶ Not every array represents a heap
- ▶ In a max-heap, the largest element is at the root and the smallest element is in a leaf

Heapify


- ▶ Assume that node i violates the heap property, but the children nodes $2i$ and $2i+1$ (if exist) are heaps.

```

HEAPIFY(A,i)
  if A[2i] > A[i] or A[2i+1] > A[i] then
    if A[2i+1] > A[2i] then
      exchange A[i] and A[2i+1];
      HEAPIFY(A, 2i+1)
    else
      exchange A[i] and A[2i];
      HEAPIFY(A, 2i)
  end
end
    
```



Heapify: example



Building a binary heap

- ▶ Given an array $A[1..n]$, build a binary heap for array elements

BUILD-HEAP(A,n)

 for $i = \lfloor n/2 \rfloor$ **downto** 1 **do** HEAPIFY(A,i);

- ▶ *Exercise:* build the heap for $A = [4, 1, 3, 2, 16, 9, 10, 14, 8, 7]$

BUILD-HEAP: complexity

- ▶ Straightforward estimation $O(n \cdot \log(n))$
- ▶ Refined analysis:
 - ▶ Cost of a call to HEAPIFY at a node depends on the height, h , of the node – $O(h)$.
 - ▶ Height of most nodes smaller $\lfloor \log(n) \rfloor$
 - ▶ Height of nodes h ranges from 0 to $\lfloor \log(n) \rfloor$
 - ▶ number of nodes of height h is at most $\lceil n / 2^{h+1} \rceil$?

Heap Characteristics

- ▶ Height = $\lfloor \log n \rfloor$
- ▶ Number of leaves = $\lceil n/2 \rceil$
- ▶ Number of nodes of height $h \leq \lceil n/2^{h+1} \rceil$
- ▶ *Proof by induction:*
 - ▶ remove all leaves from the heap
 - ▶ there remains $n - \lceil n/2 \rceil = \lfloor n/2 \rfloor$ nodes
 - ▶ the height of each node is decremented by 1
 - ▶ nb of nodes of height $h-1$ is (by induction) $\lfloor \lfloor n/2 \rfloor / 2^h \rfloor \leq \lceil n/2^{h+1} \rceil$

Tighter bound for BUILD-HEAP: $O(n)$

time of BUILD-HEAP is $\sum_{h=0}^{\lfloor \log n \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil O(h) = O\left(n \sum_{h=0}^{\lfloor \log n \rfloor} \frac{h}{2^h}\right)$

note that $\sum_{h=0}^{\lfloor \log n \rfloor} \frac{h}{2^h} \leq \sum_{h=0}^{\infty} \frac{h}{2^h} = \frac{1/2}{(1-1/2)^2} = 2$

therefore the time is $O(n)$

Priority Queue

- ▶ MAX(A): return the heap root
- ▶ EXTRACT-MAX(A):

Priority Queue


- ▶ MAX(A): return the heap root
- ▶ EXTRACT-MAX(A): exchange $A[1]$ and $A[n]$, discard element n , and apply HEAPIFY(A,1)
- ▶ INCREASE-KEY(A,i,k):

Priority Queue

- ▶ **MAX(A)**: return the heap root
- ▶ **EXTRACT-MAX(A)**: exchange $A[1]$ and $A[n]$, discard element n , and apply **HEAPIFY(A,1)**
- ▶ **INCREASE-KEY(A,i,k)**:
 $A[i] \leftarrow k$;
while $A[\lfloor i/2 \rfloor] < A[i]$ **do**
 exchange $A[\lfloor i/2 \rfloor]$ and $A[i]$;
 $i \leftarrow \text{parent}(i)$
end
- ▶ **INSERT(A,i)**: insert a new leaf $n+1$ with key $-\infty$; call **INCREASE-KEY(A,n+1,k)**

INCREASE-KEY: example

INCREASE-KEY(A,9,15)



Priority Queues: time bounds

- ▶ **MAX**: $O(1)$
- ▶ **EXTRACT-MAX**, **INCREASE-KEY**, **INSERT**: $O(\log(n))$

Priority Queues: time bounds

- ▶ **MAX**: $O(1)$
- ▶ **EXTRACT-MAX**, **INCREASE-KEY**, **INSERT**: $O(\log(n))$

Various improvements have been proposed

- ▶ *Fibonacci heaps* take $O(1)$ amortized time for **INSERT** and **INCREASE-KEY**
- ▶ if keys are integers bounded by C , *van Emde Boas trees* support **INSERT**, **DELETE**, **MAX**, **MIN**, **SUCC**, **PRED** in time $O(\log \log(C))$

Back to Dijkstra's algorithm

With adjacency matrix
time $O(n^2)$

With adjacency lists

Q : priority queue
if implemented by binary heaps:
 n building a heap of n elements: $O(n)$
 n operations **MIN**: $O(n \cdot \log n)$
 m operations **RELAX**: $O(m \cdot \log n)$
total time $O((n+m) \cdot \log n)$:
improves over $O(n^2)$ if $m = o(n^2 / \log n)$

time can be improved to $O(n \cdot \log n + m)$ using *Fibonacci heaps*,
as decreasing the key takes $O(1)$ amortized